

Moving Target Defense for the Placement of Intrusion Detection Systems in the Cloud

Sailik Sengupta¹, Ankur Chowdhary², Dijiang Huang², and
Subbarao Kambhampati¹

¹ Yochan Lab, Arizona State University, USA

² Secure Network and Computing Lab, Arizona State University, USA
{sailiks, achaud16, dijiang, rao}@asu.edu

Abstract. A lot of software systems are deployed in the cloud. Owing to realistic demands for an early product launch, oftentimes there are vulnerabilities that are present in these deployed systems (or eventually found out). The cloud service provider can find and leverage this knowledge about known vulnerabilities and the underlying communication network topology of the system to position network and host based Intrusion Detection Systems (IDS) that can effectively detect attacks. Unfortunately, deploying IDS on each host and network interface impacts the performance of the overall system. Thus, in this paper, we address the problem of placing a limited number of IDS by using the concept of Moving Target Defense (MTD). In essence, we propose an MTD system that allows a defender to shift the detection surfaces and strategically switch among the different IDS placement configurations in each round. To find a secure switching strategy, we (1) formulate the problem of placing a limited number of IDS systems in a large cloud network as a Stackelberg Game between the cloud administrator and an (external or stealthy) attacker, (2) design scalable methods to find the optimal strategies for switching IDS placements at the start of each round, and (3) formally define the problem of identifying the most critical vulnerability that should be fixed, and propose a solution for it. We compare the strategy generated by our method to other state-of-the-art strategies, showcasing the effectiveness and scalability of our method for real-world scenarios.

Keywords: Moving Target Defense · Intrusion Detection Systems · Stackelberg Games.

1 Introduction

System Administrators, oftentimes, use Intrusion Detection Systems (IDS) to detect on-going attacks on modern-day cyber-systems [34]. These IDS systems perform sophisticated operations—like signature-matching [3], anomaly detection [15, 11], machine learning [17, 1, 21] etc. — to investigate either live traffic on the wire, (using Network-based IDS (NIDS) [33, 2]) or monitor resources on a machine (using Host-based IDS (HIDS) [41, 13]) to flag anomalous request that

might result in potential loss of confidentiality, integrity or availability. Cloud service providers, who host multiple third parties on their platform, encounter non-trivial challenges when it comes to deploying these IDS that can identify vulnerabilities that exist in their system on account of legacy or operational constraints [12]. The foremost among these challenges is the placement of IDS on all nodes of a large network, which results in reduced performance [20, 40]. Moreover, third party users of the cloud platform, due to privacy and security reasons, have constraints about sharing their data with the cloud provider [6].

Thus, given a cloud service provider’s performance constraints and their customer’s privacy constraints, we look at the problem of placing a limited number of IDS systems in the various nodes of the cloud system. It is trivial to see that if we place IDS systems statically, that monitor only certain attacks on specific nodes, an attacker (especially a *stealthy one*, i.e. one who resides inside a deployed systems and can attack a node anywhere in the network as opposed to having access to only hosts at the entry point) will eventually figure out our placement strategy [40]. At this point, a strategic attacker can always select attacks that circumvent the IDS placed, thus passing through our cloud network undetected. To address this, we design a Moving Target Defense (MTD) approach for dynamic placement of IDS systems on cloud systems.

The placement mechanism for our cloud framework places both Network and Host-based IDS. We will use a NIDS called **snort** [32] and an HIDS known as **auditd** for detecting malicious behavior over the network or on a host in our cloud system. The assumption is that NIDS is placed at the gateway of each tenant network and the HIDS is deployed on each individual VM. A dynamic switching (or MTD) strategy selectively turns **on/off** the different HIDS or NIDS systems that can be used to monitor requests or hosts, thereby shifting the detection surface at each round without the need to consider switching costs among configurations because **on/off** commands from a central server sent out only at the start of every round hardly impact performance.

The cyber-security community has mostly defined and used MTD, so far, to shift the attack surface of a system that takes away the advantage of reconnaissance that an attacker has [43]. In this work, we generalize this notion of MTD and introduce an MTD system that shifts the detection surface to keep an attacker guessing about whether their next attack will be detected or not. In conjunction to that, the key contributions of this paper are,

- We formulate the problem of placing limited IDS systems in a large cloud-based network using MTD as a two-player Stackelberg Game between the defender and an attacker. The equilibrium of this game gives us the optimal movement strategy that the defender should use to switch between the various IDS placements.
- We obtain the utility values of the players in this game by combining (1) the Common Vulnerability Scoring System (CVSS) that has been previously used to represent the impact of attacks on the defender’s system [23] and (2) the centrality values of the nodes in which an IDS is deployed that lets us

- capture (i) the connectivity information and (ii) the impact on performance when an IDS is placed on that node [40].
- We design a scalable optimization problem to find the Stackelberg Equilibrium of our formulated game (Sec. 4). In this approach, we introduce an input parameter α that lets the defender balance between the security of the system and the impact on the performance of the system.
- We define the problem of finding the most critical vulnerability in an cloud environment with a strategic attacker and a multi-objective utility function and propose a method to solve it (Sec. 5).
- We demonstrate the effectiveness of our approach on a running example by comparing it to state-of-the-art deterministic, uniformly random and centrality based MTD switching strategies. We then provide experimental results in a real-world large-scale cloud-based environment that proves the scalability of our approach (Sec. 6).

2 Related Work

Moving Target Defense [43] has been recently used to thwart a wide range of attacks against network-based [16, 39] and cloud-based systems [7, 9]. These methods mostly shift the attack surfaces that takes away the advantage of reconnaissance an attacker has. A stealthy and strategic adversary [5], who can reside deep within the network, can render these methods ineffective.

For such cases, researchers have previously investigated the placement of detection systems in large network-based environments and designed both static [20] and dynamic [40] placement mechanisms based on graph-theoretic measures. Unfortunately, the former method cannot adapt its placement strategy when facing a stealthy adversary. On the other hand, the latter method, which does not incorporate the knowledge of known vulnerabilities, performs sub-optimally when facing a strategic and rational adversary.

A switching strategy for any dynamic placement method or MTD system needs to incorporate attacker modeling and thus, game theoretic reasoning for it to be effective [31, 36, 35]. Previously, authors in [22] have modeled an MTD system as a game called PLADD, based on FlipIt [38]. This work assumes that different agents control the server in different game rounds, which is an impractical setting for cloud environments. In [19], researchers assume known vulnerabilities and design a deception mechanism using a Stackelberg Model to introduce honeynets against a specific class of attackers. Authors in [35] formulate the switching between various web-stack configurations as a Stackelberg Game. Unfortunately, the methods to find the Stackelberg equilibrium in these cases become intractable as the number of defender strategies explodes combinatorially.

Researchers have shown that decomposition of the reward structure makes the problem of finding the Stackelberg Equilibrium computationally efficient [24]. We leverage this information and design the rewards for our game, while ensuring that the Stackelberg equilibrium balances between two important metrics [23]– (1) the costs of placing IDSs (on performance, cost of countermeasure deployment etc.) and (2) the impacts on the security of our system.




Fig. 1. Defender's system on the enterprise cloud that the attacker wants to attack.

Lastly, researchers have leveraged the attack graph information of a network and tried to come up with classical AI planning approaches [26] or MDP-style approaches [14, 29] to find effective ways of finding critical attacks against a system. Unfortunately, these approaches cannot be trivially applied in the case of dynamic systems like MTD and thus we develop an approach to find the most critical vulnerability that should be fixed in our system.

3 Game-Theoretic Modeling

In this section, we first define the threat model of our system, defining the players, their action/strategy sets using a small real-world scenario that we set up on an enterprise cloud (Fig. 1). We then describe how the rewards of this game are formulated leveraging the CVSS data and network topology information.

Threat Model In our attack model, we consider a multi-tenant cloud network. The controller node, shown in the Fig. 1, is used for network management and orchestration. The network administrator (or the defender) utilizes a management network to access controller nodes and cloud servers hosting VMs. We consider two agents– the defender \mathcal{D} , who is trying to deploy IDS and an (external or stealthy) attacker \mathcal{A} , who is trying to remain undetected while attacking the system. As a running example, we will use the scenario deployed by \mathcal{D} shown in Fig 1. Furthermore, this system has a set of known vulnerabilities, that are yet to be fixed and as per our assumptions, known to both the agents \mathcal{D} and \mathcal{A} .

ID	VM	c_b	Vulnerability	CVE ID	IOC
a_1	G1	4	SSH Buffer Overflow	CVE-2016-6289	NIDS sshAlert
a_2	G2	7	rlogin	CVE-1999-0651	NIDS rlogin
a_3	W	0	Cross Side Scripting	CVE-2016-2163	HIDS webAccess
a_4	D	0	Weak Credentials	CVE-2001-0839	HIDS fileIntegrity
a_5	F	0	vsftpd backdoor	CVE-2015-1419	HIDS ftpLogin

Table 1. The different VMs in the defender’s network, their betweenness centrality (c_b) in the graph, the known vulnerabilities in these nodes (VMs), and the corresponding Network/Host-based Intrusion Detection Systems (NIDS/HIDS) which can detect these attacks, also known as the Indicators of Compromise (IOC).

We assume that the attacker \mathcal{A} can be located either inside or outside the cloud network. The attacker’s primary goal is to (1) compromise a VM using known vulnerabilities and remain (2) undetected while doing so. Since, the attacker can utilize network probing to identify the OS and software versions, it will eventually get to know the vulnerabilities (CVEs) associated with the system, and can then systematically exploit these in order to obtain network access or elevated privileges. Furthermore, the attacker can only be detected when it attacks a vulnerability for which the corresponding IDS is in place at the time of exploitation. For stealthy attackers [5], who have spend a lot of cost and/or effort in gaining access to an internal node, the latter is of utmost importance.

Now given the system’s communication graph, we extract the set A of all the n known vulnerabilities in our system ($n = (|A|)$). For our system, we choose represent an *attack* (and the corresponding IDS that detects this attack) using the ID in the first column of table 1. Thus, $n = 5$ and the set $A = \{a_1, a_2, a_3, a_4, a_5\}$. Note that this ID encodes a two-tuple $\langle \text{MachineName}, \text{CVE-ID} \rangle$. Thus, multiple attacks corresponding to a single machine will each receive a unique ID.

The defender \mathcal{D} , as mentioned before, has a limited budget to place only $k (< n)$ IDS mechanisms due to resource constraints. Also, we assume that, due to privacy constraints, \mathcal{D} cannot place an IDS mechanism on the ‘SQL Server (M)’ (shown in Figure 1). Thus, in our model, we disregard any vulnerabilities present on this node. (Note that our model can detect vulnerabilities that trigger NIDS alarms on the network interface G1 and affect M. But we exclude such vulnerabilities from our example.) . Now, \mathcal{D} has $\binom{n}{k}$ ways in which it can deploy the k IDSs. This is the action set of \mathcal{D} . Formally, the defender’s action set is denoted by the set $A_k = \{S \in A : |S| = k\}$. In the running example, we will assume that $k = 2$. Thus, the defender’s action set is:

$$\{(a_1, a_2), (a_1, a_3), (a_1, a_4), (a_1, a_5), (a_2, a_3), (a_2, a_4), (a_2, a_5), (a_3, a_4), (a_3, a_5), (a_4, a_5)\}$$

Since, we assume a strong adversary who either knows or can find out all the attacks in our system, the action set of the attacker is the attack set $A = \{a_1, a_2, a_3, a_4, a_5\}$ itself.

In game theory, this action set is often referred to as the set of pure strategies, where each action (either a placement strategy or an attack) is a pure strategy (for \mathcal{D} or \mathcal{A} respectively). As stated earlier, if a defender chooses a pure strategy, i.e, any one out of the ten pure strategies shown, to deploy k IDS systems, the

attacker, with reconnaissance on its side, will eventually figure out \mathcal{D} 's strategy and start choosing attacks that do not trigger these alarms. In order to address this limitation, the defender can play a mixed strategy, i.e. have a probability associated with playing each pure strategy and at the start of each round picks one by randomly sampling a pure strategy from the set of pure strategies. Note that this is similar to applying the concept of Moving Target Defense where the defender chooses to switch randomly among the different deployment configurations (i.e. by choosing one of the ten IDS placements in our case) at the start of each time period.

Common Vulnerability Scoring System (CVSS) The CVSS metric provides two quantitative scores for each CVE present in our system—(1) the Impact Score (IS) that represents the effect a particular attack has on the Confidentiality, Integrity, and Availability of a system and (2) the Exploitability Score (ES), which encodes the complexity of actually exploiting a particular vulnerability. The system defines a way to combine both of these scores to calculate a third score, known as the Base Score (BS) that tries to consider both the impact of an attack *vs.* the difficulty in exploiting it.

The CVSS scores thus leverage the knowledge of cyber security experts across the globe to provide a numerical value corresponding to each (known) vulnerability that reflect its severity and expertise necessary to exploit it. We, inspired by other research works before us [42, 27, 35], use the CVSS to calibrate the reward values of our game.

3.1 Stackelberg Games

Having defined the players and their action (or pure strategy) sets, there are additional real-world aspects that we want to incorporate in the formulation of our game. One such aspect is that the defender, who hosts the system that an attacker attacks, plays first. To accurately model this fact, we use the concept of Stackelberg models in which one player (\mathcal{D}) is able to act (or commit to act) before the other player (\mathcal{A}) plays. In such games, the well known concept of Nash Equilibrium gives sub-optimal utilities to the players and thus one needs to find the Stackelberg Equilibrium of these games, in which the leader's (\mathcal{D}) strategy is contingent upon the fact that the follower (\mathcal{A}) can observe \mathcal{D} 's strategy and play accordingly. Thus, in this adversarial leader-follower game, \mathcal{D} can simulate \mathcal{A} in their mind and decide on a mixed strategy that gives it the highest utility keeping in mind (that a rational) \mathcal{A} will choose the best action ($\in A$), i.e. the action that maximizes \mathcal{A} 's reward, in response.

3.2 Utility Modeling

Having designed the action sets of both the players, we can now specify the utilities for both the players when each of them commits to a pure strategy. Note that just to enumerate all the utility values for our game we would have

to specify $2 \cdot \binom{n}{k} \cdot n$ values corresponding to the reward/utility values for each of the players \mathcal{D} and \mathcal{A} in the normal form game matrix. With this general reward structure, finding the mixed-strategy Stackelberg equilibrium of this game would be computationally inefficient, specifically $O(\binom{n}{k})$ [10]. Thus, we now devise a particular reward structure that captures all the aspects of our problem and lets us efficiently compute the equilibrium strategy.

For each attack $a \in A$, if \mathcal{D} places an IDS to detect it, we will say that \mathcal{D} *covers* it. Otherwise, we say that a is left *uncovered*. Since the defender has resources to only deploy k IDSs, k elements of A will be covered and the remaining $n - k$ attacks will remain uncovered at any point in time. We will now decompose the reward structure of this game and define four types of utility values corresponding to each attack $a \in A$.

$$\langle U_{c,a}^{\mathcal{D}}, U_{u,a}^{\mathcal{D}}, U_{c,a}^{\mathcal{A}}, U_{u,a}^{\mathcal{A}} \rangle$$

where $U_{c,a}^{\mathcal{D}}$ and $U_{u,a}^{\mathcal{D}}$ denotes the utilities a defender gets for covering and not covering an attack a respectively. Similarly, $U_{c,a}^{\mathcal{A}}$ and $U_{u,a}^{\mathcal{A}}$ represent the utility the attacker gets when they use an attack a that is covered (and thus gets detected) and not covered (and thus avoids detection) respectively. The values for these symbols are obtained by leveraging the knowledge of security experts encoded in the Common Vulnerabilities Scoring System (CVSS) [28] and the realistic costs of deploying IDSs. For each attack a in our model, we will represent these scores as IS_a , ES_a and BS_a using CVSS metrics, previously discussed in Sec. 3.

Cost of deploying IDS. We denote the cost of deploying an IDS corresponding to an attack $a \in A$ as \hat{c}_a . For our example, we assume the cost of deploying an IDS (shown in the IOC column of table 1) to be proportional to the betweenness centrality of the VMs on which the IDS is deployed because a VM with high betweenness centrality will affect the latency of routing packets or the latency of processing a request. Also, the centrality values are normalized in the interval $[0, 10]$ to be comparable to the CVSS metrics IS_a , ES_a and BS_a 3. Note that the model in this paper allows another user to define \hat{c}_a in a different way.

We now leverage these defined metrics to design the following rewards for the four utilities our each attack a present in our system,

$$\begin{aligned} U_{c,a}^{\mathcal{D}} &= -1 * \hat{c}_a, U_{u,a}^{\mathcal{D}} = -1 * IS_a \\ U_{c,a}^{\mathcal{A}} &= -1 * ES_a, U_{u,a}^{\mathcal{A}} = +1 * BS_a \end{aligned}$$

We now provide some rationale for modeling the rewards in this particular manner. The value of $U_{c,a}^{\mathcal{D}}$ is negative since even if it detected an attack, it incurred a cost in order to detect it and moreover there is no extra positive reward given to \mathcal{D} for protecting their system, which is supposed to be the primary functionality. When \mathcal{D} does not place an IDS for detecting the attack a , it incurs a negative utility ($U_{u,a}^{\mathcal{D}}$ equal to IS_a if the attacker uses attack a).

For the attacker \mathcal{A} , if it chooses an attack action a which the defender covers (i.e. can detect), it gets a negative utility $U_{c,a}^{\mathcal{A}}$ proportional to the time and cost it

had to invest in doing it, which is (somewhat) measured by ES . Also, as \mathcal{A} gains nothing by doing this attack (since the defender can deploy a countermeasure on detection [8]), no positive value is added to it. Lastly, when the attacker uses an attack for which the defender has not placed an IDS, we give a positive utility that (conceptually) adds the IS and subtracts the cost (ES) of performing the attack. Since BS already captures this trade-off, we use it directly.

4 Computing the Stackelberg Equilibrium

We need to solve for the Stackelberg Equilibrium of our game to obtain probability values for each configuration mentioned in A_k , where $A_k \subset A$ such that $|A_k| = k$. Unfortunately, since there are $\binom{n}{k}$ such probabilities (corresponding to each element in A_k), solving for all these variables at once will not yield an efficient solution. Instead, we will solve for the probabilities p_a which represents the probability that a certain attack $a \in A$ is covered by an IDS in a round.

To that extent, we first describe a method that can help in generating the marginal strategies for the defender by solving n ($= |A|$) Linear Programs. Note that this is a polynomial time solution and only possible because of the particular reward structure we use in our game. Then, we shall propose an efficient Mixed Integer Quadratic Program (MIQP) method based on this method that helps us to obtain the same marginal strategy, but by solving just one optimization problem. We show that although this formulation, in the general case, is known to be computationally hard to solve, in our case, by efficient use of the branch-and-cut mechanism, can be solved in polynomial time.

4.1 Multiple LP method

Let T denote the set of k tokens that the defender \mathcal{D} can allocate to cover k of the n attacks. Allocating a token to an attack a means that \mathcal{D} has placed the IDS that can detect the particular attack. Now, let the variables p_a represent the probability with which an attack a is covered by one of the k tokens and $p_{a,t}$ represent the probability with which a particular attack a is covered by a particular token $t \in T$. Having defined the probabilities p_a , the defender's expected utility for deploying an IDS to detect a particular attack a^* should be $U_{c,a^*}^D * p_{a^*} + U_{c,a^*}^D * (1 - p_{a^*})$ [24, 25]. Note that, for our scenario, this does not capture the cost \mathcal{D} incurs in deploying the other $k - 1$ IDS mechanisms. Thus, we modify the defender's utility to $U_{c,a^*}^D * (1 - p_{a^*}) + \frac{1}{k} \sum_{a \in A} U_{c,a}^D * p_a$, where the second term denotes the average cost for a particular deployment configuration.

On the other hand, we can simply define the attacker's expected utility for using a particular attack a as $U_{c,a}^A * p_a + U_{c,a}^A * (1 - p_a)$. We now present the optimization problem that maximizes the defender's objective function and the

attacker's utility given that an attacker chooses to use the attack a^* .

$$\begin{aligned}
 \max \quad & \alpha \cdot \frac{1}{k} \sum_{a \in A} U_{c,a}^{\mathcal{D}} p_a + (1 - \alpha) \cdot U_{u,a^*}^{\mathcal{D}} (1 - p_{a^*}) \\
 \text{s.t.} \quad & p_a \in [0, 1] \quad \forall a \in A \\
 & p_{t,a} \in [0, 1] \quad \forall a \in A, t \in T \\
 & \sum_{a \in A} p_{t,a} = 1 \quad \forall t \in T \\
 & \sum_{t \in T} p_{t,a} = p_a \quad \forall a \in A \\
 & U_{c,a}^{\mathcal{A}} p_a + U_{u,a}^{\mathcal{A}} (1 - p_a) \leq U_{c,a^*}^{\mathcal{A}} p_{a^*} + U_{u,a^*}^{\mathcal{A}} (1 - p_{a^*})
 \end{aligned} \tag{1}$$

where α is an input parameter that allows the defender to trade the performance of the system with respect to the security of the system (and vice versa). In the extreme case when $\alpha = 0$, the defender optimizes only for security and completely ignores the fact that deploying k IDSs might affect the performance of the system. In this case, as shown in 6, \mathcal{D} lands up randomizing more between the deployment configurations of the system. On the other hand, when $\alpha = 1$, the defender optimizes for performance, hardly placing an IDS on systems that affect performance even when it is detrimental to security. We discuss the effects of selecting various α -s in section 6.

Before we dive into what the constraints mean, note that this is a Linear Program (LP) and thus, can be solved in polynomial time. The first two sets of constraints ensure that the optimization variables p_a and $p_{t,a}$ are valid probabilities. The third set of constraints ensures that every token is fully utilized in covering the different attacks in A . The equality of this constraint is possible in our case since (1) all our tokens are homogeneous, i.e. any token $t \in T$ can be used to cover any attack $a \in A$ and (2) the number of tokens k ($= |T|$) is less than the number of attacks n ($= |A|$). Thus, we prune away solutions that do not fully utilize all the tokens. The fourth set of constraints ensure that the probabilities of allocating various tokens to cover an attack a add up to the probability that a is covered. The final set of constraints ensure that the attacker selecting a^* maximizes their utility. Lastly, note that given the values of $p_{t,a}$ one can easily obtain p_a using the fourth set of constraints.

To obtain the (globally) optimal solution (and thus find the optimal marginal strategy) for the defender, we can iterate over all the n attack choices made by the attacker and pick the solution that maximizes \mathcal{D} 's utility. Note that, here we enforce the attacker to select a pure strategy as opposed to a mixed strategy. This is not a limitation since for any mixed strategy the attacker can pick in this Stackelberg Game, there always exists a pure strategy in support of it [30].

As the number of VMs and vulnerabilities, i.e., n , increase, this solution method needs to solve a large number of LPs. Thus, We now propose an efficient MIQP that solves for the solution at one go and provide an efficient branch-and-cut algorithm for solving it in polynomial.

4.2 Compiling Multiple LPs into an Efficient Mixed Integer Quadratic Program (MIQP)

Now, we first introduce n binary switch variables, one for each attack $a \in A$ and represent it as w_a . When the attacker exploits vulnerability a (i.e. uses the attack action a), $w_a = 1$. Otherwise, $w_a = 0$. We now propose the following optimization problem,

$$\begin{aligned}
\max \quad & \alpha \cdot \frac{1}{k} \sum_{a \in A} U_{c,a}^{\mathcal{D}} p_a + (1 - \alpha) \cdot w_a * U_{u,a}^{\mathcal{D}} (1 - p_a) \quad (2) \\
s.t. \quad & w_a \in \{0, 1\} \quad \forall a \in A \\
& p_a \in [0, 1] \quad \forall a \in A \\
& p_{t,a} \in [0, 1] \quad \forall a \in A, t \in T \\
& \sum_{a \in A} w_a = 1 \\
& \sum_{a \in A} p_{t,a} = 1 \quad \forall t \in T \\
& \sum_{t \in T} p_{t,a} = p_a \quad \forall a \in A \\
& 0 \leq v_a - (U_{c,a}^{\mathcal{A}} p_a + U_{u,a}^{\mathcal{A}} (1 - p_a)) \leq (1 - w_a) * M \quad \forall a \in A
\end{aligned}$$

where M represents a large number with respect to the maximum reward the attacker can get, i.e. $M \gg 10$, and v_a is the utility value of the attacker at equilibrium. The first constraint ensures that the switch variables are binary. The fourth constraint enforces the attacker to select a pure strategy since the switch variable corresponding to only one attack can be turned. As mentioned in the previous section, this is not a limiting assumption. Lastly, the final set of constraints encodes the complementary slackness condition of the attacker's utility maximization problem [30].

As the defender plays first, it can reason about the attacker picking each attack and select the strategy which gives \mathcal{D} the maximum reward. If the attacker responds to the defender's strategy with attack a^* , then $w_{a^*} = 1$. In that case, the RHS of the last constraint (with a^*) becomes zero and along with the LHS, equality holds. Thus, v_{a^*} is \mathcal{A} 's utility value. For all the other attacks $a (\neq a^*)$ that were not selected by \mathcal{A} , both the inequalities can be trivially satisfied (as M is a large number) by selecting an appropriate value for v_a .

Theorem 1. *MIQP defined in equation 2 produces the same solution as the set of LPs described in equation 1.*

Proof. Let us say that when attacker selects an attack a^1 , the defender gets the highest utility as per equation 1. Now, let us say that equation 2 decides that the defender's utility is strictly better when attacker selects any another attack $a^2 (\neq a^1)$, and thus, $w_{a^2} = 1$. Notice that if this is true, then the objective function value of LP when $a^* = a^2$ is strictly greater than the objective function




Fig. 2. Branch-and-cut tree for the proposed MIQP.

Attack	a_1	a_2	a_3	a_4	a_5
$U_{c,a}^D$	-5.7	-10.0	0.0	0.0	0.0
$U_{u,a}^D$	-6.4	-6.4	-2.9	-6.4	-2.9
$U_{c,a}^A$	-8.6	-10	-8.6	-10	-10
$U_{u,a}^A$	6.8	7.5	4.3	7.5	5.0

Table 2. Player utilities for each vulnerability depending on whether an IDS is deployed to detect any attacks that exploit it.

	a_1	a_2	a_3	a_4	a_5
t_1	0	0.44	0	0.22	0.34
t_2	0.45	0	0.34	0.21	0

Table 3. Probability of allocating a token (in order to deploy the corresponding IDS) for detecting each attack.

value of the LP with $a^* = a^1$. But that is a contradiction. Hence, the MIQP defined in equation 2 must select a for the attacker.

Similarly, we can prove the other way—that a solution that is optimal for the MIQP (eqn. 2) is also optimal for the LP case. ■

Theorem 2. *MIQP defined in equation 2 can be solved in polynomial time with the branch-and-cut method.*

Proof. To prove this, we first represent the branch-and-cut tree for our MIQP in fig. 2. In that, notice that the right children (shown in red) correspond to an LP problem (similar to the one defined in equation 1) where only a particular attack a_i is selected ($w_{a_i} = 1$) and other attacks are not used by the attacker. Since no children of any right child (red node) can generate another solution, the search tree below them can be pruned away. Now, the tree can have at most $n - 1$ left children which corresponds to at most n right children, which in turn corresponds to at most n LP problems that need to be solved. Since each LP can be solved in polynomial time and we will solve no more than n LPs, this MIQP can be solved in polynomial time. ■

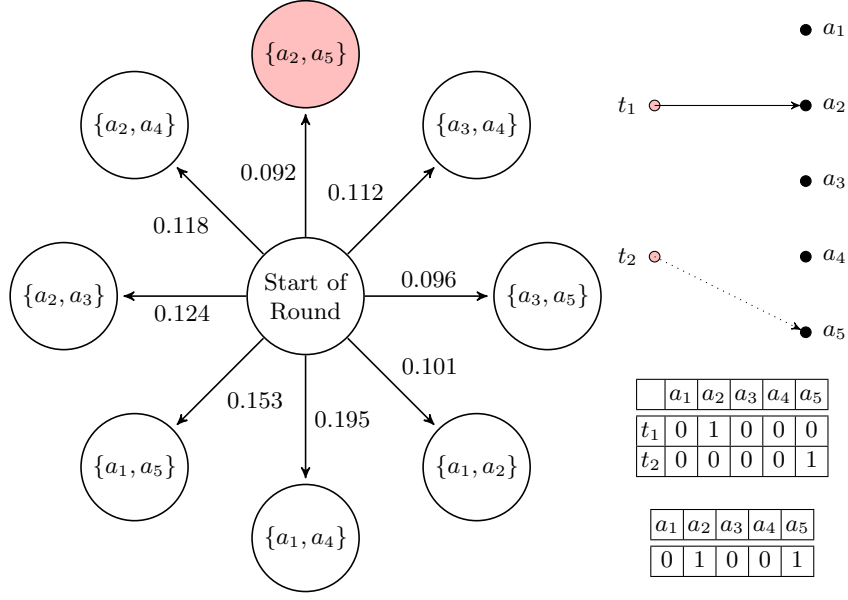


Fig. 3. Optimal mixed strategy of the defender for our scenario (when $\alpha = 0.1$). The probability values for picking up one of the eight IDS placements at the start of each round are written on the edges. For the strategy $\{a_2, a_5\}$ (colored in Pink), the allocation matrix is shown on the right.

4.3 Obtaining Implementable Strategies

Although we have obtained the values p_a and $p_{t,a}$, there are no guarantees that we will be able to convert these marginal probabilities into $\binom{n}{k}$ probability values that correspond to a defender's deployment strategies, i.e. one that can be implemented in practice. In order to convert these into *implementable strategies*, we use the general version of the Birkhoff Von-Neumann Theorem as stated in [25]. We state this here for completeness.

Birkhoff Von-Neumann Theorem. Consider an $k \times n$ matrix P with real numbers $p_{t,a} \in [0, 1]$, such that for each $1 \leq t \leq k$, $\sum_{a=1}^n p_{t,a} \leq 1$, and for each $1 \leq a \leq n$, $\sum_{t=1}^k p_{t,a} \leq 1$. Then, there exist matrices P^1, P^2, \dots, P^q and weights $w^1, w^2, \dots, w^q \in (0, 1]$, such that (1) $\sum_{x=1}^q w^x = 1$; (2) $\sum_{x=1}^q w^x P^x = M$; (3) for each $1 \leq x \leq q$, the elements of M^x are $p_{t,a}^x \in \{0, 1\}$ and (4) for each $1 \leq x \leq q$, we have for each $1 \leq t \leq k$, $\sum_{a=1}^n p_{t,a}^x \leq 1$ and for each $1 \leq a \leq n$, $\sum_{t=1}^k p_{t,a}^x \leq 1$.

This theorem guarantees that given the probability matrix $p_{t,a}$, we can always obtain the $\binom{n}{k}$ implementable probabilities. The third and fourth equalities in the optimization problem in 1 ensure that the constraint structure imposed on P is a *bihierarchy*, which authors in [4] show as a sufficient condition for any marginal probability matrix P to be *implementable*.

For our example, assuming that the cost associated with deploying each IDS on a certain VM is a function of the latency it creates. Furthermore, since VMs

Input: Utility Matrix
Output: a^*
Result: Finds and outputs the most critical vulnerability that results in the highest defender utility when fixed
 $\text{max_def_util} \leftarrow -\infty$;
 $a^* \leftarrow \text{None}$;
while $a \in A$ **do**
 $A' \leftarrow A \setminus a$;
 $\text{obj_val}, \leftarrow \text{solve MIQP (2) with action set } A'$;
 if $\text{obj_val} > \text{max_def_util}$ **then**
 $\text{max_def_util} \leftarrow \text{obj_val}$;
 $a^* \leftarrow a$;
 end
 return a^*
end

Algorithm 1: Algorithm to find the most critical vulnerability in the Defender's system, which when fixed results in the highest utility.

that are responsible for communication between other VMs would impact the latency the most when an IDS is placed on it. Thus, we assume time impact on the overall latency of the system is proportional to the betweenness centrality of the nodes in our network (normalized between $[0, 10]$). With that, the utility values for the attacker and defender are shown in table 2. We first use these values to solve for the optimal marginal strategy (shown in Fig. 3) using the MIQP described in 2. We then use Theorem 1 to obtain the mixed strategies that the defender can actually use to deploy the IDS systems (shown in Fig. 3).

5 Identifying the Most Critical Vulnerability

In real-world scenarios, system administrators, who have a list of known vulnerabilities it should address, have limited developer resources to fix all of the known CVEs in their system at once. Thus, the question of which vulnerability they should fix in order to improve the security of the system is a critical one. In our case, since (1) the rewards of the formulated game are not zero-sum and (2) the defender wants to balance a multi-objective function (that tries to balance the security and usability metrics), figuring out the (critical) vulnerability that \mathcal{D} needs to fix becomes even more difficult.

Given that we can find the utilities for the defender using equation 2, we can ask the question *which attack a when removed would produce the maximum utility for \mathcal{D}* . A simple algorithm would be to iterate over all the attacks, removing them one by one, reformulating the MIQP and selecting the attack that maximizes the defender's utility when removed. We describe this idea formally in algorithm 1 and use it to find the most critical vulnerability of our system. The utilities obtained by removing one vulnerability at a time are shown below (for $\alpha = 0.1$).

$$\langle a_1 : -1.90; a_2 : -1.70; a_3 : -2.30; a_4 : -2.23; a_5 : -2.27 \rangle$$

Thus, in our system, a_2 is the most critical vulnerability since fixing a_2 will result in the highest (gain in) defender’s utility.

6 Experiments

We present the results of two different experiments– (1) comparison of our placement strategy (Fig. 3) with existing approaches, and (2) implementation of the Stackelberg Game Strategy (SGS) on a large cloud network instance.

6.1 Comparison with Existing Strategies

In this section, we compare our approach to three other different MTD strategies in the context of our running example where $n = 5$ and $k = 2$ –

(1) *Deterministic Pure Strategy (DPS)*. This strategy selects a single pure strategy out of the $\binom{5}{2}$ placement strategies. As per work by [20], for DPS, we place IDS to detect a_1 and a_2 (since $G1$ and $G2$ are the most critical VMs), which are on the critical paths for any attack flow. Note that, in the context of a stealthy attacker who can exploit any vulnerability in the system, definition of a critical node, on which an IDS can be deployed, is not clear. Thus, DPS has an inherent disadvantage when compared to MTD strategies, which we now describe.

(2) *Uniform Random Strategy (URS)*. In this case, we select each of the $\binom{5}{2}$ placements or pure strategies with an equal probability of 0.1. In this case, each attack a is covered in four (out of the ten) pure strategies since having placed an IDS (or token which denotes an IDS was placed) for a , there are $\binom{4}{1} = 4$ ways of placing the other token. Thus, the marginal probabilities are $0.1 * 4 = 0.4$.

(3) *Centrality Based Strategy (CBS)*. This strategy, motivated in the work by [40], has previously been shown to be effective for detecting stealthy botnets when PageRank is used as a centrality measure. Since our network is an undirected graph, we use the betweenness centrality measure for evaluation. Since only two of our nodes ($G1$ and $G2$) have non-zero values for betweenness centrality, we switch between seven of the ten configurations– three in which only a_1 is covered, three in which only a_2 is covered and one in which both a_1 and a_2 are covered. Since $G1$, on which a_1 is present has a lower centrality value in comparison to $G2$, on which a_2 is present, the first three configurations are less likely than the next three. The last configuration, in which both a_1 and a_2 are covered, is the most likely configuration. The marginal probabilities for covering each attack in the system, as per this strategy, is shown in Fig. 4.

	a_1	a_2	a_3	a_4	a_5
URS	0.4	0.4	0.4	0.4	0.4
DPS	1	1	0	0	0
CBS	0.52	0.73	0.25	0.25	0.25

Fig. 4. Table showcasing the marginal probabilities with which IDS is placed on a node for the different strategies.

Effectiveness of Our Approach We plot the defender’s utility value for our approach and compare it to all the other approaches. The results are shown in Fig. 5. When adversaries are strategic, i.e. can reason about defender strategies




Fig. 5. Defender's utility for the various MTD strategies as the security-usability trade-off value (α) varies from zero to one.

and act rationally to maximize their utility, our method clearly dominates the other methods (see the plots for CBS(min), URS(min) and DPS). On the other hand, if \mathcal{A} is completely irrational, i.e., selects attacks in a way to maximize \mathcal{D} profit, we consider the best case for the other MTD strategies (see URS(max) and CBS(max)), it turns out that only URS is a little better when $\alpha \in (0.5, 0.97]$. In this range, our algorithm selects nodes with high centrality measure to improve security in the case of a strategic attacker. This increases the deployment cost and reduces the multi-objective function value, letting URS dominate. CBS on the other hand with no information about the known attacks or performance costs, switches only among the useless and performance expensive configurations, being strictly dominated by SGS. Note that all the mechanisms we compare against here do not adapt to the security and performance trade-off that is important to the defender. Thus, as the value of α changes, the marginal probabilities of CBS, URS and DPS being constant, result in the plot being a straight line. On the other hand, SGS, our intelligent switching mechanism, solves the multi-objective optimization when coming up with its mixed strategy.

When α is low (i.e. $\in [0, 0.3]$), our method switches among eight out of the ten pure strategies. As α increases further and the costs start to matter, it places IDS systems more on nodes that impact performance the least. Beyond a certain value (when $\alpha > 0.7$) it realizes that the cost of placing IDS on G1 and G2 (for detecting $a1$ and $a2$) are extremely high on performance of the system and sticks to only (three) strategies where both of these are not covered.

6.2 Testing on a Large Cloud Network

We conduct a test on a Mininet [37] based network simulator. The setup comprised of 15 VMs and 42 CVEs distributed uniformly on a 15 VM flat network 10.0.0.0/24, as shown in the Fig. 6. We measure the throughput for the server (10.0.0.15) hosting an ssh application on port 5002. We used Snort NIDS [32] pre-configured with the attack signatures corresponding to the CVEs reported




Fig. 6. Flat network with 15 VMs hosting multiple NIDS and HIDS

by the vulnerability scanner. For instance, the attack signature below checks the payload for shellcode targeting remote buffer overflow vulnerability on ssh service running on port 5022.

```

alert tcp any any -> 10.0.0.15 5022 (msg:"EXPLOIT ssh remote
overflow"; content:"/bin/sh"; reference:bugtraq,2347;
reference:cve,2008-5161; sid:1324; rev:6;)

```

Network-based IDS Snort was configured to run in IDS (intrusion-detection) as well as IPS (intrusion-prevention) mode. The AF Packet, which is an IPS configuration, creates a bridge between inspected interfaces (e.g., h1-eth1:s1-eth1). This leads to increased packet processing latency since each packet on a particular bridge is inspected against all traffic patterns which are part of signatures.

Host-based IDS auditd [18] was configured to monitor file integrity of configuration files such as `/etc/sshd.conf` and binary files for vulnerable services present on the network. A daemon was configured on each inspected host to generate an alert if there is a change in the hash value of inspected files.

The goal of this experiment was to measure the impact of the HIDS/NIDS deployment on the throughput of the services being accessed by normal users. We show that as \mathcal{D} places more IDSs (1 to 15), we observe a substantial drop in the throughput of the system from 18 Gbps to 6 Gbps (see Fig. 7). This shows that deployment of NIDS and HIDS without considering the impact on network latency can affect the Quality of Service (QoS) for legitimate users in a cloud network.

In Figure 8, we vary the number of IDS systems placed in the system and see how the defender utilities vary. Initially, as the number of IDS increases from 2 to 17, the defender’s utility increases at a slow rate since there are too few IDS systems to detect attacks on all the 42 vulnerabilities. As the number of IDS systems are increased beyond 18, the defender’s utility starts to increase substantially in each step. At this point if the attacker does not pick their attack strategically, it is detected with high probability. However placement of more IDSs beyond a certain point (30) as shown in the Fig. 8, results in a substantial decrease in throughput, outweighing the benefits of security provided by IDS. Lastly, The most critical vulnerability found in this system was CVE-2013-2207.




Fig. 7. Throughput of the MiniNET network vs the number of IDS systems deployed on the system increase.




Fig. 8. Defender's utility value with increase in the number of IDS systems placed on the system.

7 Conclusion and Future Work

In this paper, we addressed the problem of placing a fixed number of IDS systems in a large cloud environment by proposing a Moving Target Defense (MTD) approach for shifting the detection surface. We formulated this problem as a two-player general-sum Stackelberg Game between the cloud administrator (our defender) and an attacker and designed two scalable algorithms that can (1) find the Stackelberg Equilibrium of the formulated game, which lets the cloud service provided a balance between the security and usability of their system, and (2) find the most critical vulnerability in their system. We assumed that the attacker is rational, i.e. he will try to exploit the known vulnerabilities present of VMs in the cloud network by scanning the network. A sophisticated attacker may however perform reconnaissance over an extended period of time and may thus go undetected by IDS. We plan to model a sophisticated attacker as a future work for this paper.

References

1. Al-Jarrah, O., Arafat, A.: Network intrusion detection system using neural network classification of attack behavior. *Journal of Advances in Information Technology* Vol **6**(1) (2015)
2. Bakshi, A., Dujodwala, Y.B.: Securing cloud from ddos attacks using intrusion detection system in virtual machine. In: *Communication Software and Networks*, 2010. ICCSN'10. Second International Conference on. pp. 260–264. IEEE (2010)
3. Brown, D.J., Suckow, B., Wang, T.: A survey of intrusion detection systems. Department of Computer Science, University of California, San Diego (2002)
4. Budish, E., Che, Y.K., Kojima, F., Milgrom, P.: Designing random allocation mechanisms: Theory and applications. *American Economic Review* **103**(2), 585–623 (2013)
5. Center, M.I.: Apt1: Exposing one of chinas cyber espionage units. Mandian. com (2013)

6. Chen, D., Zhao, H.: Data security and privacy protection issues in cloud computing. In: Computer Science and Electronics Engineering (ICCSEE), 2012 International Conference on. vol. 1, pp. 647–651. IEEE (2012)
7. Chowdhary, A., Alshamrani, A., Huang, D., Liang, H.: Mtd analysis and evaluation framework in software defined network (mason). In: Proceedings of the 2018 ACM International Workshop on Security in Software Defined Networks & Network Function Virtualization. pp. 43–48. ACM (2018)
8. Chowdhary, A., Pisharody, S., Huang, D.: Sdn based scalable mtd solution in cloud network. In: Proceedings of the 2016 ACM Workshop on Moving Target Defense. pp. 27–36. ACM (2016)
9. Chung, C.J., Khatkar, P., Xing, T., Lee, J., Huang, D.: Nice: Network intrusion detection and countermeasure selection in virtual network systems. *IEEE transactions on dependable and secure computing* **10**(4), 198–211 (2013)
10. Conitzer, V., Sandholm, T.: Computing the optimal strategy to commit to. In: Proceedings of the 7th ACM conference on Electronic commerce. pp. 82–90. ACM (2006)
11. Dastjerdi, A.V., Bakar, K.A., Tabatabaei, S.G.H.: Distributed intrusion detection in clouds using mobile agents. In: Advanced Engineering Computing and Applications in Sciences, 2009. ADVCOMP'09. Third International Conference on. pp. 175–180. IEEE (2009)
12. Debar, H., Dacier, M., Wespi, A.: Towards a taxonomy of intrusion-detection systems. *Computer Networks* **31**(8), 805–822 (1999)
13. Deshpande, P., Sharma, S., Peddoju, S., Junaid, S.: Hids: A host based intrusion detection system for cloud computing environment. *International Journal of System Assurance Engineering and Management* pp. 1–10 (2014)
14. Durkota, K., Lisý, V., Bosanský, B., Kiekintveld, C.: Optimal network security hardening using attack graph games. In: IJCAI. pp. 526–532 (2015)
15. Garfinkel, T., Rosenblum, M., et al.: A virtual machine introspection based architecture for intrusion detection. In: *Ndss*. vol. 3, pp. 191–206 (2003)
16. Hu, Z., Zhu, M., Liu, P.: Online algorithms for adaptive cyber defense on bayesian attack graphs (2017)
17. Ibrahim, L.M.: Anomaly network intrusion detection system based on distributed time-delay neural network (dtdnn). *Journal of Engineering Science and Technology* **5**(4), 457–471 (2010)
18. Ilgun, K., USTAT, A.: A real-time intrusion detection system for unix. University of California Santa Barbara Master Thesis (1992)
19. Jajodia, S., Park, N., Serra, E., Subrahmanian, V.: Share: A stackelberg honey-based adversarial reasoning engine. *ACM Transactions on Internet Technology (TOIT)* **18**(3), 30 (2018)
20. Jha, S., Sheyner, O., Wing, J.: Two formal analyses of attack graphs. In: Computer Security Foundations Workshop, 2002. Proceedings. 15th IEEE. pp. 49–63. IEEE (2002)
21. Jo, S., Sung, H., Ahn, B.: A comparative study on the performance of intrusion detection using decision tree and artificial neural network models. *Journal of the Korea Society of Digital Industry and Information Management* **11**(4), 33–45 (2015)
22. Jones, S., Outkin, A., Gearhart, J., Hobbs, J., Siirola, J., Phillips, C., Verzi, S., Tauritz, D., Mulder, S., Naugle, A.: Evaluating moving target defense with pladd. Tech. rep., Sandia National Labs-NM, Albuquerque (2015)
23. Kheir, N., Cuppens-Boulahia, N., Cuppens, F., Debar, H.: A service dependency model for cost-sensitive intrusion response. In: European Symposium on Research in Computer Security. pp. 626–642. Springer (2010)

24. Kiekintveld, C., Jain, M., Tsai, J., Pita, J., Ordóñez, F., Tambe, M.: Computing optimal randomized resource allocations for massive security games. In: Proceedings of The 8th International Conference on Autonomous Agents and Multiagent Systems-Volume 1. pp. 689–696. International Foundation for Autonomous Agents and Multiagent Systems (2009)
25. Korzhyk, D., Conitzer, V., Parr, R.: Complexity of computing optimal stackelberg strategies in security resource allocation games. In: AAAI (2010)
26. Letchford, J., Vorobeychik, Y.: Optimal interdiction of attack plans. In: Proceedings of the 2013 international conference on Autonomous agents and multi-agent systems. pp. 199–206. International Foundation for Autonomous Agents and Multiagent Systems (2013)
27. Maleki, H., Valizadeh, S., Koch, W., Bestavros, A., van Dijk, M.: Markov modeling of moving target defense games. In: Proceedings of the 2016 ACM Workshop on Moving Target Defense. pp. 81–92. ACM (2016)
28. Mell, P., Scarfone, K., Romanosky, S.: Common vulnerability scoring system. IEEE Security & Privacy 4(6) (2006)
29. Panda, S., Vorobeychik, Y.: Near-optimal interdiction of factored mdps. In: Conference on Uncertainty in Artificial Intelligence (2017)
30. Paruchuri, P., Pearce, J.P., Marecki, J., Tambe, M., Ordóñez, F., Kraus, S.: Playing games for security: An efficient exact algorithm for solving bayesian stackelberg games. In: Proceedings of the 7th international joint conference on Autonomous agents and multiagent systems-Volume 2. pp. 895–902. International Foundation for Autonomous Agents and Multiagent Systems (2008)
31. Pita, J., Jain, M., Marecki, J., Ordóñez, F., Portway, C., Tambe, M., Western, C., Paruchuri, P., Kraus, S.: Deployed armor protection: the application of a game theoretic model for security at the los angeles international airport. In: Proceedings of the 7th international joint conference on Autonomous agents and multiagent systems: industrial track. pp. 125–132. AAMAS (2008)
32. Roesch, M., et al.: Snort: Lightweight intrusion detection for networks. In: Lisa. vol. 99, pp. 229–238 (1999)
33. Roschke, S., Cheng, F., Meinel, C.: An extensible and virtualization-compatible ids management architecture. In: Information Assurance and Security, 2009. IAS’09. Fifth International Conference on. vol. 2, pp. 130–134. IEEE (2009)
34. Rowland, C.H.: Intrusion detection system (Jun 11 2002), uS Patent 6,405,318
35. Sengupta, S., Vadlamudi, S.G., Kambhampati, S., Doupé, A., Zhao, Z., Taguinod, M., Ahn, G.J.: A game theoretic approach to strategy generation for moving target defense in web applications. In: Proceedings of the 16th Conference on Autonomous Agents and MultiAgent Systems. pp. 178–186. International Foundation for Autonomous Agents and Multiagent Systems (2017)
36. Sinha, A., Nguyen, T.H., Kar, D., Brown, M., Tambe, M., Jiang, A.X.: From physical security to cybersecurity. Journal of Cybersecurity 1(1), 19–35 (2015)
37. Team, M.: Mininet (2014)
38. Van Dijk, M., Juels, A., Oprea, A., Rivest, R.L.: Flipit: The game of “stealthy takeover”. Journal of Cryptology 26(4), 655–713 (2013)
39. Venkatesan, S., Albanese, M., Amin, K., Jajodia, S., Wright, M.: A moving target defense approach to mitigate ddos attacks against proxy-based architectures. In: Communications and Network Security (CNS), 2016 IEEE Conference on. pp. 198–206. IEEE (2016)
40. Venkatesan, S., Albanese, M., Cybenko, G., Jajodia, S.: A moving target defense approach to disrupting stealthy botnets. In: Proceedings of the 2016 ACM Workshop on Moving Target Defense. pp. 37–46. ACM (2016)

41. Vieira, K., Schulter, A., Westphall, C., Westphall, C.: Intrusion detection for grid and cloud computing. *It Professional* **12**(4), 38–43 (2010)
42. Zhu, Q., Başar, T.: Game-theoretic approach to feedback-driven multi-stage moving target defense. In: *International Conference on Decision and Game Theory for Security*. pp. 246–263. Springer (2013)
43. Zhuang, R., DeLoach, S.A., Ou, X.: Towards a theory of moving target defense. In: *Proceedings of the First ACM Workshop on Moving Target Defense*. pp. 31–40. ACM (2014)